

Assignment 11

File: WC_Mapper.java

```
1.      package com.javatpoint;
2.
3.      import java.io.IOException;
4.      import java.util.StringTokenizer;
5.      import org.apache.hadoop.io.IntWritable;
6.      import org.apache.hadoop.io.LongWritable;
7.      import org.apache.hadoop.io.Text;
8.      import org.apache.hadoop.mapred.MapReduceBase;
9.      import org.apache.hadoop.mapred.Mapper;
10.     import org.apache.hadoop.mapred.OutputCollector;
11.     import org.apache.hadoop.mapred.Reporter;
12.     public class WC_Mapper extends MapReduceBase implements Mapper<Long
        Writable,Text,Text,IntWritable>{
13.         private final static IntWritable one = new IntWritable(1);
14.         private Text word = new Text();
15.         public void map(LongWritable key, Text value,OutputCollector<Text,IntW
            ritable> output,
16.             Reporter reporter) throws IOException{
17.             String line = value.toString();
18.             StringTokenizer tokenizer = new StringTokenizer(line);
19.             while (tokenizer.hasMoreTokens()){
20.                 word.set(tokenizer.nextToken());
21.                 output.collect(word, one);
22.             }
23.         }
24.
25.     }
```

File: WC_Reducer.java

```

1.      package com.javatpoint;
2.
3.      import java.io.IOException;
4.      import java.util.Iterator;
5.      import org.apache.hadoop.io.IntWritable;
6.      import org.apache.hadoop.io.Text;
7.      import org.apache.hadoop.mapred.MapReduceBase;
8.      import org.apache.hadoop.mapred.OutputCollector;
9.      import org.apache.hadoop.mapred.Reducer;
10.
11.     public class WC_Reducer extends MapReduceBase implements Reducer<
12.         Text,IntWritable,Text,IntWritable> {
13.
14.         public void reduce(Text key, Iterator<IntWritable> values,OutputCollector
15.             <Text,IntWritable> output,
16.             Reporter reporter) throws IOException {
17.             int sum=0;
18.             while (values.hasNext()) {
19.                 sum+=values.next().get();
20.             }
21.             output.collect(key,new IntWritable(sum));
22.         }
23.     }

```

```

1.      package com.javatpoint;
2.
3.      import java.io.IOException;
4.      import org.apache.hadoop.fs.Path;
5.      import org.apache.hadoop.io.IntWritable;
6.      import org.apache.hadoop.io.Text;
7.      import org.apache.hadoop.mapred.FileInputFormat;
8.      import org.apache.hadoop.mapred.FileOutputFormat;

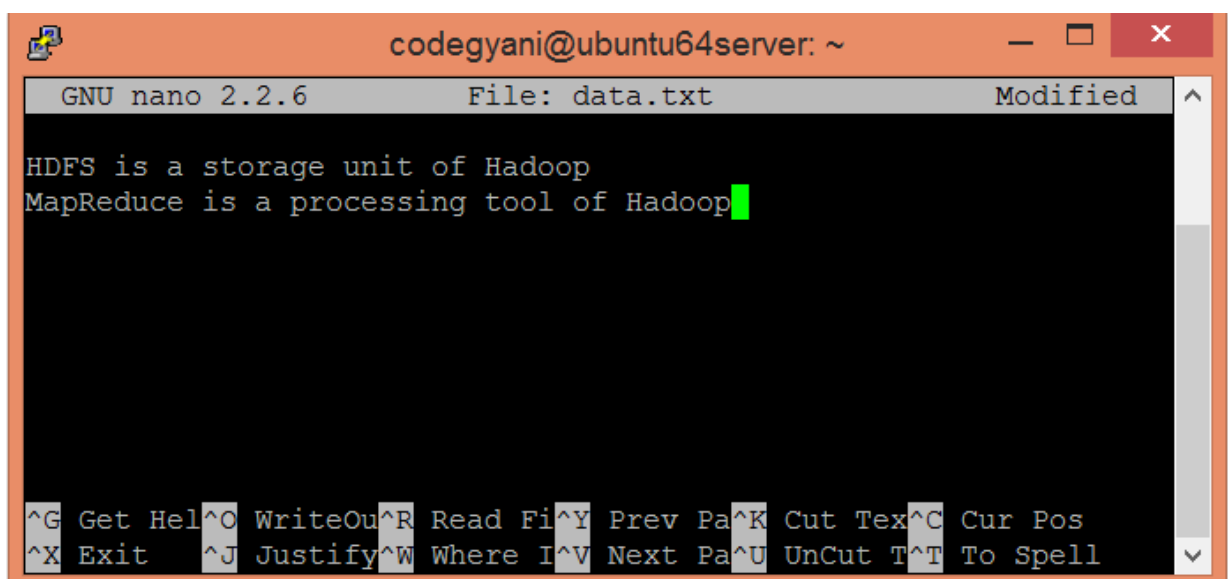
```

```

9.      import org.apache.hadoop.mapred.JobClient;
10.     import org.apache.hadoop.mapred.JobConf;
11.     import org.apache.hadoop.mapred.TextInputFormat;
12.     import org.apache.hadoop.mapred.TextOutputFormat;
13.     public class WC_Runner {
14.         public static void main(String[] args) throws IOException{
15.             JobConf conf = new JobConf(WC_Runner.class);
16.             conf.setJobName("WordCount");
17.             conf.setOutputKeyClass(Text.class);
18.             conf.setOutputValueClass(IntWritable.class);
19.             conf.setMapperClass(WC_Mapper.class);
20.             conf.setCombinerClass(WC_Reducer.class);
21.             conf.setReducerClass(WC_Reducer.class);
22.             conf.setInputFormat(TextInputFormat.class);
23.             conf.setOutputFormat(TextOutputFormat.class);
24.             FileInputFormat.setInputPaths(conf,new Path(args[0]));
25.             FileOutputFormat.setOutputPath(conf,new Path(args[1]));
26.             JobClient.runJob(conf);
27.         }
28.     }

```

Output:



```

codegyani@ubuntu64server: ~
GNU nano 2.2.6      File: data.txt      Modified
HDFS is a storage unit of Hadoop
MapReduce is a processing tool of Hadoop
^G Get Hel^O WriteOu^R Read Fi^Y Prev Pa^K Cut Tex^C Cur Pos
^X Exit    ^J Justify^W Where I^V Next Pa^U UnCut T^T To Spell

```

```
codegyani@ubuntu64server: ~
codegyani@ubuntu64server:~$ nano data.txt
codegyani@ubuntu64server:~$ cat data.txt
HDFS is a storage unit of Hadoop
MapReduce is a processing tool of Hadoop
codegyani@ubuntu64server:~$
```

Hadoop Overview Datanodes Snapshot Startup Progress Utilities ▾

Browse Directory

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	codegyani	supergroup	74 B	2/11/2019, 3:12:12 PM	1	128 MB	data.txt

Hadoop, 2015.

Browse Directory

/test

Go!

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	codegyani	supergroup	74 B	2/11/2019, 3:12:12 PM	1	128 MB	data.txt

Hadoop, 2015.

Hadoop

Overview

Datanodes

Snapshot

Startup Progress

Utilities

Browse Directory

/r_output

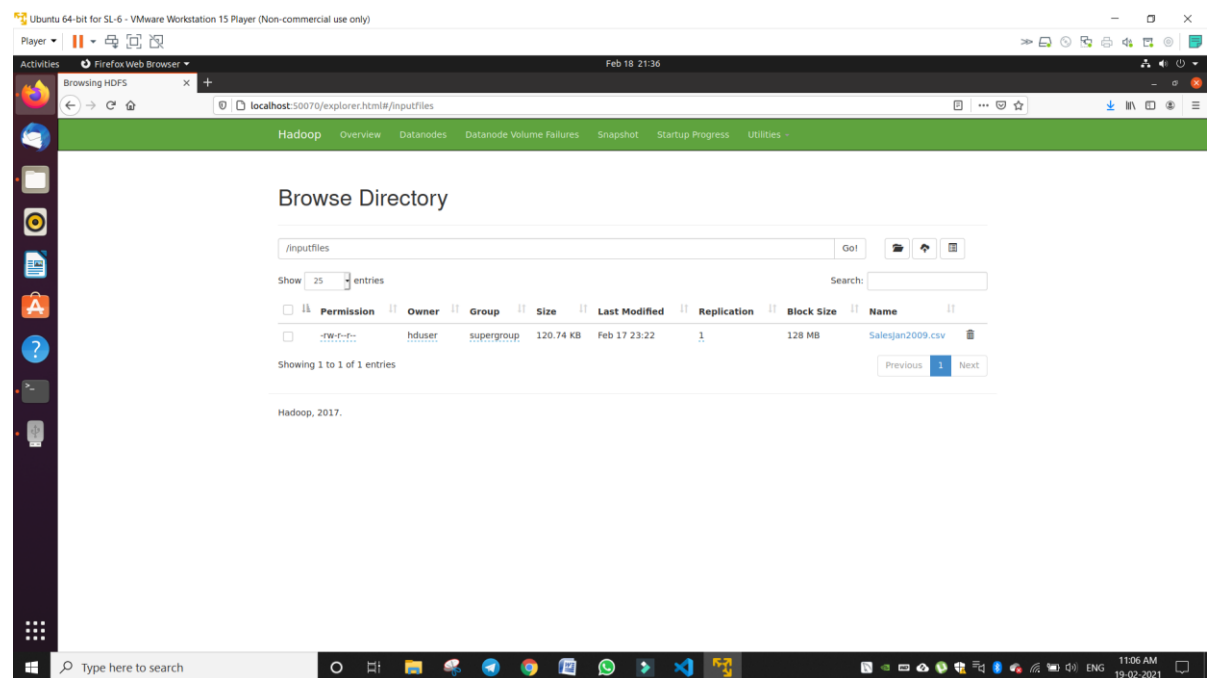
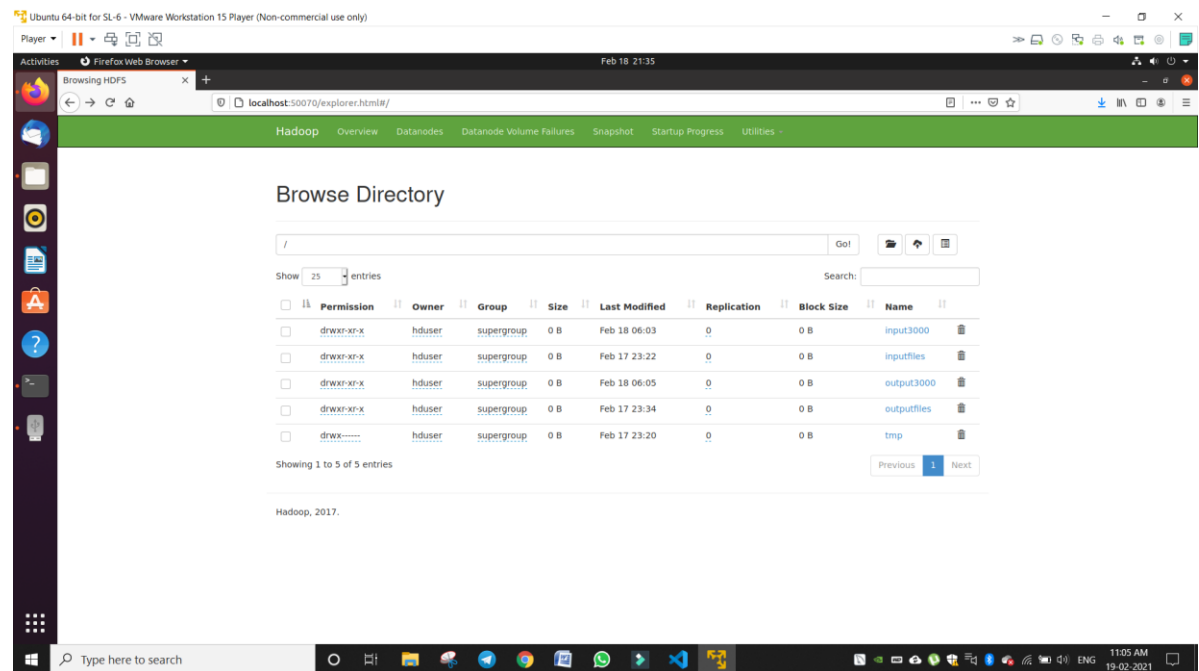
Go!

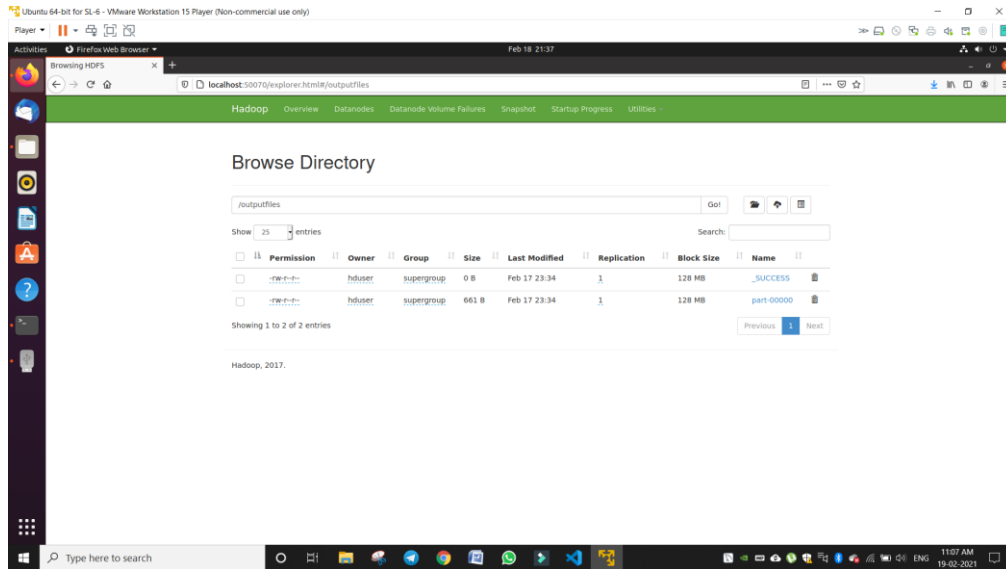
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	codegyani	supergroup	0 B	2/11/2019, 3:52:27 PM	1	128 MB	_SUCCESS
-rw-r--r--	codegyani	supergroup	79 B	2/11/2019, 3:52:23 PM	1	128 MB	part-00000

codegyani@ubuntu64server: ~

```
codegyani@ubuntu64server:~$ hdfs dfs -cat /r_output/part-00000
HDFS      1
Hadoop    2
MapReduce      1
a          2
is          2
of          2
processing    1
storage      1
tool         1
unit         1
codegyani@ubuntu64server:~$
```

Assignment 12





Assignment 13

Program :

```
package com.org.vasanth.weather;
```

```
import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;
import java.util.StringTokenizer;
```

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.KeyValueTextInputFormat;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;
```

```
/**
 *
 *
 * This is an Hadoop Map/Reduce application for Working on weather data It reads
 * the text input files, breaks each line into stations weather data and finds
 * average for temperature , dew point , wind speed. The output is a locally
 * sorted list of stations and its 12 attribute vector of average temp , dew ,
 * wind speed of 4 sections for each month.
 *
 * To run: bin/hadoop jar target/weather-1.0.jar [-m <i>maps</i>] [-r
 * <i>reduces</i>] <i>in-dir for job 1</i> <i>out-dir for job 1</i> <i>out-dir
 * for job 2</i>
 */
public class Weather extends Configured implements Tool {
    final long DEFAULT_SPLIT_SIZE = 128 * 1024 * 1024;

    /**
     * Map Class for Job 1
     *
     * For each line of input, emits key value pair with
     * station_yearmonth_sectionno as key and 3 attribute vector with
     * temperature , dew point , wind speed as value. Map method will strip the
     * day and hour from field and replace it with section no (
     * <b>station_yearmonth_sectionno</b>, <b><temperature,dew point , wind
     * speed></b>).
     */
    public static class MapClass extends MapReduceBase
```



```

implements Mapper<LongWritable, Text, Text, Text> {

private Text word = new Text();
private Text values = new Text();

public void map(LongWritable key, Text value,
                OutputCollector<Text, Text> output,
                Reporter reporter) throws IOException {
    String line = value.toString();
    StringTokenizer itr = new StringTokenizer(line);
    int counter = 0;
    String key_out = null;
    String value_str = null;
    boolean skip = false;
    loop:while (itr.hasMoreTokens() && counter<13) {
        String str = itr.nextToken();
        switch (counter) {
            case 0:
                key_out = str;
                if(str.contains("STN")){//Ignoring rows where station id is
                    skip = true;
                    break loop;
                }else{
                    break;
                }
            case 2:
                int hour =
Integer.valueOf(str.substring(str.lastIndexOf("_")+1, str.length()));
                str = str.substring(4,str.lastIndexOf("_")-2);
                /*if(hour<=5){
                    str = str.concat("_section4");
                }else if(hour>5 && hour<=11){
                    str = str.concat("_section1");
                }else if(hour>11 && hour<=17){
                    str = str.concat("_section2");
                }else if(hour>17 && hour<=23){
                    str = str.concat("_section3");
                }*/

                if(hour>4 && hour<=10){
                    str = str.concat("_section1");
                }else if(hour>10 && hour<=16){
                    str = str.concat("_section2");
                }else if(hour>16 && hour<=22){
                    str = str.concat("_section3");
                }else{
                    str = str.concat("_section4");
                }

                key_out = key_out.concat("_").concat(str);
                break;
            case 3://Temperature
                if(str.equals("9999.9")){//Ignoring rows where
                    skip = true;

```

```

                break loop;
            }else{
                value_str = str.concat(" ");
                break;
            }
        case 4://Dew point
            if(str.equals("9999.9")){//Ignoring rows where dew point
                skip = true;
                break loop;
            }else{
                value_str = value_str.concat(str).concat(" ");
                break;
            }
        case 12://Wind speed
            if(str.equals("999.9")){//Ignoring rows where wind speed
                skip = true;
                break loop;
            }else{
                value_str = value_str.concat(str).concat(" ");
                break;
            }
        default:
            break;
    }
    counter++;
}
if(!skip){
    word.set(key_out);
    values.set(value_str);
    output.collect(word, values);
}
}

/**
 * Map Class for Job 2
 *
 * For each input, emits key value pair with station_yearmonth as key and 3
 * attribute vector with temperature , dew point , wind speed as value by
 * stripping the section no from key and adding section no into vector value
 * ( <b>station_yearmonth</b>, <b><temperature,dew point , wind speed></b>).
 */
public static class MapClassForJob2 extends MapReduceBase
    implements Mapper<Text, Text, Text, Text> {
    private Text key_text = new Text();
    private Text value_text = new Text();
    @Override
    public void map(Text key, Text value,
        OutputCollector<Text, Text> output, Reporter reporter) throws
        IOException {

        String str = key.toString();
        String station = str.substring(str.lastIndexOf("_")+1, str.length());
        str = str.substring(0,str.lastIndexOf("_"));
        key_text.set(str);

```

```

        StringTokenizer itr = new StringTokenizer(value.toString());
        String str_out = station.concat("<");
        while (itr.hasMoreTokens()) {
            String nextToken = itr.nextToken(" ");
            str_out = str_out.concat(nextToken);
            str_out = ((itr.hasMoreTokens()) ? str_out.concat(",") :
str_out.concat(">"));
        }
        value_text.set(str_out);
        output.collect(key_text,value_text);
    }
}

/**
 * Reducer Class for Job 1
 *
 * A reducer class that just emits 3 attribute vector with average
 * temperature , dew point , wind speed for each of the section of the month
 * for each input
 */
public static class Reduce extends MapReduceBase
    implements Reducer<Text, Text, Text, Text> {
    private Text value_out_text = new Text();
    public void reduce(Text key, Iterator<Text> values,
        OutputCollector<Text, Text> output, Reporter reporter) throws
IOException {

        double sum_temp = 0;
        double sum_dew = 0;
        double sum_wind = 0;
        int count = 0;

        while (values.hasNext()) {

            String str = values.next().toString();

            StringTokenizer itr = new StringTokenizer(str);
            int count_vector = 0;
            while (itr.hasMoreTokens()) {
                String nextToken = itr.nextToken(" ");
                if(count_vector==0){
                    sum_temp += Double.valueOf(nextToken);
                }
                if(count_vector==1){
                    sum_dew += Double.valueOf(nextToken);
                }
                if(count_vector==2){
                    sum_wind += Double.valueOf(nextToken);
                }
                count_vector++;
            }
            count++;
        }
        double avg_tmp = sum_temp / count;
        double avg_dew = sum_dew / count;
        double avg_wind = sum_wind / count;

```

```

        System.out.println(key.toString()+" count is "+count+" sum of temp is
"+sum_temp+" sum of dew is "+sum_dew+" sum of wind is "+sum_wind+"\n");

        String value_out = String.valueOf(avg_tmp).concat("
").concat(String.valueOf(avg_dew)).concat(" ").concat(String.valueOf(avg_wind));
        value_out_text.set(value_out);
        output.collect(key, value_out_text);
    }
}

/**
 * Reducer Class for Job 2
 *
 * A reducer class that just emits 12 attribute vector with average
 * temperature , dew point , wind speed for all section of the month
 * for each input
 */
public static class ReduceForJob2 extends MapReduceBase
    implements Reducer<Text, Text, Text, Text> {
    private Text value_out_text = new Text();
    public void reduce(Text key, Iterator<Text> values,
        OutputCollector<Text, Text> output, Reporter reporter) throws
IOException {
        String value_out = "";
        while (values.hasNext()) {
            value_out = value_out.concat(values.next().toString()).concat(" ");
        }
        value_out_text.set(value_out);
        output.collect(key, value_out_text);
    }
}

static int printUsage() {
    System.out.println("weather [-m <maps>] [-r <reduces>] <job_1 input> <job_1
output> <job_2 output>");
    ToolRunner.printGenericCommandUsage(System.out);
    return -1;
}

/**
 * The main driver for weather map/reduce program.
 * Invoke this method to submit the map/reduce job.
 * @throws IOException When there is communication problems with the
 *         job tracker.
 */
public int run(String[] args) throws Exception {
    Configuration config = getConf();

    // We need to lower input block size by factor of two.
    /*config.setLong(
        FileInputFormat.SPLIT_MAXSIZE,
        config.getLong(
            FileInputFormat.SPLIT_MAXSIZE, DEFAULT_SPLIT_SIZE) / 2);*/

    JobConf conf = new JobConf(config, Weather.class);
    conf.setJobName("Weather Job1");

```

```

// the keys are words (strings)
conf.setOutputKeyClass(Text.class);
// the values are counts (ints)
conf.setOutputValueClass(Text.class);

conf.setMapOutputKeyClass(Text.class);
conf.setMapOutputValueClass(Text.class);

conf.setMapperClass(MapClass.class);
//conf.setCombinerClass(Combiner.class);
conf.setReducerClass(Reduce.class);
List<String> other_args = new ArrayList<String>();
for(int i=0; i < args.length; ++i) {
    try {
        if ("-m".equals(args[i])) {
            conf.setNumMapTasks(Integer.parseInt(args[++i]));
        } else if ("-r".equals(args[i])) {
            conf.setNumReduceTasks(Integer.parseInt(args[++i]));
        } else {
            other_args.add(args[i]);
        }
    } catch (NumberFormatException except) {
        System.out.println("ERROR: Integer expected instead of " + args[i]);
        return printUsage();
    } catch (ArrayIndexOutOfBoundsException except) {
        System.out.println("ERROR: Required parameter missing from " +
            args[i-1]);
        return printUsage();
    }
}
// Make sure there are exactly 2 parameters left.
/*if (other_args.size() != 3) {
    System.out.println("ERROR: Wrong number of parameters: " +
        other_args.size() + " instead of 3.");
    return printUsage();
}*/
FileInputFormat.setInputPaths(conf, other_args.get(0));
FileOutputFormat.setOutputPath(conf, new Path(other_args.get(1)));

JobClient.runJob(conf);

JobConf conf2 = new JobConf(config, Weather.class);
conf2.setJobName("Weather Job 2");

// the keys are words (strings)
conf2.setOutputKeyClass(Text.class);
// the values are counts (ints)
conf2.setOutputValueClass(Text.class);

conf2.setInputFormat(KeyValueTextInputFormat.class);

conf2.setMapOutputKeyClass(Text.class);
conf2.setMapOutputValueClass(Text.class);

conf2.setMapperClass(MapClassForJob2.class);
//conf.setCombinerClass(Combiner.class);

```

```
conf2.setReducerClass(ReduceForJob2.class);
```

```
FileInputFormat.setInputPaths(conf2, new Path(other_args.get(1)));  
FileOutputFormat.setOutputPath(conf2, new Path(other_args.get(2)));
```

```
JobClient.runJob(conf2);  
return 0;
```

```
}
```

```
public static void main(String[] args) throws Exception {  
    int res = ToolRunner.run(new Configuration(), new Weather(), args);  
    System.exit(res);  
}
```

```
}
```

Browsing HDFS




localhost:50070/explorer.html#/outputtemp

HadoopOverviewDatanodesDatanode Volume FailuresSnapshotStartup ProgressUtilities

Browse Directory



/outputtemp

Go!



Show 25 entries

Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
<input type="checkbox"/>	-rw-r--r--	hduser	supergroup	0 B	Jun 07 22:54	1	128 MB	_SUCCESS 
<input type="checkbox"/>	-rw-r--r--	hduser	supergroup	0 B	Jun 07 22:54	1	128 MB	part-r-00000 

Showing 1 to 2 of 2 entries

Previous1Next

Hadoop, 2017.